

# Nicole Chiou

[nicchiou@gmail.com](mailto:nicchiou@gmail.com) | [linkedin.com/in/nicole-chiou](https://linkedin.com/in/nicole-chiou)

## ACADEMIC HISTORY

<b>Ph.D. Student in Computer Science</b> <i>Stanford University</i> , Stanford, CA <i>Cumulative GPA:</i> 4.02/4.00	Sep. 2022 – Present
<b>Master of Science in Robotics</b> <i>University of Pennsylvania</i> , Philadelphia, PA <i>Cumulative GPA:</i> 4.00/4.00 <i>Thesis:</i> Cross-Domain Transfer Learning with Auxiliary Task Selection for Cardiac Arrhythmia Classification <i>Advisors:</i> Dr. Rahul Mangharam and Dr. Eric Eaton	May 2020 – May 2021
<b>Bachelor of Science in Bioengineering, Minor in Computer Science</b> <i>University of Pennsylvania</i> , Philadelphia, PA <i>Cumulative GPA:</i> 3.51/4.00, <i>Last two years GPA:</i> 3.92/4.00	Aug. 2016 – May 2020

## RESEARCH INTERESTS

Causal inference, distribution shift, measurement error, data-centric approaches, trustworthy machine learning, fairness, applications in medicine and healthcare.

## RELEVANT COURSEWORK

Foundations of Causal Machine Learning, Transfer Learning, Algorithmic Fairness, Machine Learning.

## RESEARCH EXPERIENCE

<b>Machine Learning Research Assistant</b> , STAIR Lab <i>Advisor: Dr. Sanmi Koyejo, Department of Computer Science</i> <i>Stanford University</i>	May 2022 – Present
<ul style="list-style-type: none"><li>Developing dataset sub-selection heuristics for multi-source domain accumulation (data addition) settings. Adapting a causal modeling perspective to characterize and quantify distribution shifts between source domains and an observed target domain.</li><li>Establishing self-supervised distillation methods that leverage synthetic data to train LM-as-judge models to assess whether LM-generated medical texts are factually consistent with inputs. Designing automated risk categorization metrics to enable the evaluation of LMs in making deployment safety decisions.</li><li>Defining mechanisms to combine DSPy-based prompt optimization with a counterfactual-based prompting strategy to simulate varying covariates of health equity. Quantifying model bias and potential downstream harms of medical chatbots.</li><li>Operationalizing the evaluation of fairness in multi-stage decision-making pipelines. Decomposing LM-as-judge pipelines into auditable components (e.g., validator, generator). Determining necessary and sufficient conditions to bridge the gap between the global fairness objective and local fairness in intermediate stages of the pipeline.</li><li>Investigating the role of standard risk scoring measures in health outcome inequalities using causal mediation analysis to quantify unexplained inequalities that persist even after equalizing for environmental and procedural variables.</li></ul>	
<b>Machine Learning Research Assistant</b> , KoyejoLab <i>Advisor: Dr. Sanmi Koyejo, Department of Computer Science</i> <i>University of Illinois Urbana-Champaign</i>	Aug. 2021 – May 2022

<b>Machine Learning Research Assistant</b> , mLAB <i>Advisor: Dr. Rahul Mangharam, Departments of Electrical Engineering and Computer Science</i> <i>University of Pennsylvania</i>	Jan. 2020 – Aug. 2021
<ul style="list-style-type: none"><li>Performed patient-specific optimization of implantable cardioverter-defibrillator (ICD) algorithms using data-driven methods and gradient-free optimization of parametric methods.</li></ul>	

- Designed generative modeling-based approaches for domain adaptation, deriving and incorporating a latent context embedding from each patient's normal rhythm to address the challenge of limited patient-specific data.
- Established a neural network-based discriminator for cardiac arrhythmias using transfer learning and multi-task learning to leverage learned representations from electrocardiogram (ECG) data, improving classification performance on class-imbalanced patient-specific electrogram (EGM) data.
- Developed a multi-task learning objective and a technique for auxiliary task selection for hierarchical and class-imbalanced data.

## WORK EXPERIENCE

---

<b>Google Student Researcher, Ph.D.</b> , Google Research, Responsible AI <i>Dr. Mercy Asiedu and Dr. Katherine Heller</i>	May 2023 – Nov. 2023
<ul style="list-style-type: none"> <li>• Developed deep learning methods to reduce the subjectivity of predicting fetal hypoxia from visual cardiotocography (CTG) interpretation.</li> <li>• Analyzed the effect of (a) the choice of objective vs. subjective ground truth labels and (b) temporal distribution shift on predictive performance.</li> <li>• Proposed end-to-end preprocessing, data augmentation, and statistical evaluation methods to overcome challenges with the limited dataset size.</li> </ul>	
<b>Google Student Researcher, Ph.D.</b> , Google Brain (now Google DeepMind)	May 2022 – Nov. 2022
<ul style="list-style-type: none"> <li>• Estimated the optimal target predictor in an unsupervised domain adaptation setting, where the source differs from the target due to a shift in the distribution of the latent subgroup.</li> <li>• Showed that causally-informed latent variable models can estimate the distribution of the latent subgroup variable with the help of concept and proxy variables available in the source domain.</li> <li>• Evaluated the approach on a medical classification task with demographic subpopulation shift.</li> </ul>	
<b>Bio1 Systems</b> , San Carlos, California <i>Machine Learning Consultant, Software</i>	May 2021 – Oct. 2021
<ul style="list-style-type: none"> <li>• Phase II proposal contributor for the Army Applied SBIR Program; selected for funding in October 2021.</li> <li>• Developed an ML clinical decision support tool that combines raw sensor and EHR data to provide medical alerts.</li> <li>• Designed recurrent neural network (RNN) approaches to detect exposure to a chemical or biological agent-based on heart rate data obtained from wearable devices and predict the mortality risk for affected patients based on dynamic vital sign data and cross-sectional EHR data.</li> <li>• Established an end-to-end training framework, from data integration and cleaning to evaluation and on-device deployment.</li> </ul>	

## TEACHING EXPERIENCE

---

<b>Applied Machine Learning</b> (CS 441), taught by Dr. Marco Morales	Fall 2021
<b>Artificial Intelligence</b> (CIS 521), taught by Dr. Chris Callison-Burch	Summer 2021
<b>Computing Systems</b> (CIT 593), taught by Dr. Thomas Farmer	Summer 2020
<b>Biological Data Science</b> (ENM 375), taught by Dr. Jennifer Cremins	Spring 2019

## MEMBERSHIPS

---

<b>Women in Machine Learning (WiML)</b> , Mentor	2022 – Present
<b>Penn Engineering Master's Advisory Board (EMAB)</b> , Robotics representative	2020 – 2021
<b>Women in Computer Science (WiCS)</b> , Mentor	2018 – 2021

## HONORS & AWARDS

---

<b>HAI-HPI Research Fellow</b> , Human-Centered AI	Jan. 2024 – Present
<b>Outstanding Teaching Assistant Award</b> , Computing Systems	Aug. 2020
<b>Bioengineering Senior Design Award</b> , relievVR Bioreponsive Virtual Reality	May 2020
<b>2nd Place Johns Hopkins International Healthcare Design Competition (200+ teams)</b> , relievVR	Apr. 2020
<b>Penn Berkman Innovation Prize Winner</b> , relievVR	Oct. 2019

## PUBLICATIONS

---

\* denotes equal contribution.  $\alpha\text{-}\beta$  denotes alphabetical order.

### Are Domain Generalization Benchmarks with Accuracy on the Line Misspecified?

Olawale Salaudeen, **Nicole Chiou\***, Shiny Weng\*, Sanmi Koyejo.

*Transactions on Machine Learning Research (TMLR)*, 2025. [Featured in Stanford HAI News](#). Awarded a TMLR Journal-to-Conference Certification.

### On Domain Generalization Datasets as Proxy Benchmarks for Causal Representation Learning

Olawale Salaudeen, **Nicole Chiou**, Sanmi Koyejo.

*NeurIPS Causal Representation Learning Workshop (oral)*, 2024.

### Development and Evaluation of Deep Learning Models for Cardiotocography Interpretation

**Nicole Chiou**, Nichole Young-Lin, Christopher Kelly, Julie Cattiau, Tiya Tiyasirichokchai, Abdoulaye Diack, Sanmi Koyejo, Katherine Heller, Mercy Asiedu.

*npj Women's Health*, 2025.

### Bridging Gaps in Automated Acute Myocardial Infarction Detection between High-income and Low-income Countries

**Nicole Chiou**, Sanmi Koyejo, Christine Ngaruiya.

*PLOS Global Public Health*, 2024.

### Proxy Methods for Domain Adaptation

Katherine Tsai, Stephen R Pfohl, Olawale Salaudeen, **Nicole Chiou**, Matt Kusner, Alexander D'Amour, Sanmi Koyejo, Arthur Gretton.

*International Conference on Artificial Intelligence and Statistics (AISTATS)*, 2024.

### Adapting to Latent Subgroup Shifts via Concepts and Proxies

$\alpha\text{-}\beta$ . Ibrahim Alabdulmohsin\*, **Nicole Chiou\***, Alexander D'Amour\*, Arthur Gretton\*, Sanmi Koyejo\*, Matt Kusner, Stephen Pfohl\*, Olawale Salaudeen\*, Jessica Schrouff\*, Katherine Tsai\*.

*The International Conference on Artificial Intelligence and Statistics (AISTATS)*, 2023.

### Single-Trial Detection and Classification of Event-Related Optical Signals for a Brain–Computer Interface Application

**Nicole Chiou**, Mehmet Günal, Sanmi Koyejo, David Perpetuini, Antonio Maria Chiarelli, Kathy A Low, Monica Fabiani, Gabriele Gratton.

*Bioengineering*, 2024.

### Fast Optical Signals for Real-Time Retinotopy and Brain Computer Interface

David Perpetuini, Mehmet Günal, **Nicole Chiou**, Sanmi Koyejo, Kyle Mathewson, Kathy A Low, Monica Fabiani, Gabriele Gratton, Antonio Maria Chiarelli.

*Bioengineering*, 2023.

### Cross-Domain Transfer Learning with Auxiliary Task Selection for Cardiac Arrhythmia Classification

**Nicole Chiou**, Eric Eaton, Rahul Mangharam.

*University of Pennsylvania*, 2021.

## WORKS IN PROGRESS

---

### MedVAL: Toward Expert-Level Medical Text Validation with Language Models

Asad Aali, Vasiliki Bikia, Maya Varma, **Nicole Chiou**, Sophie Ostmeier, Arnav Singhvi, Magdalini Paschali, Ashwin Kumar, Andrew Johnston, Karimar Amador-Martinez, Eduardo Juan Perez Guerrero, Paola Naovi Cruz Rivera, Sergios Gatidis, Christian Bluethgen, Eduardo Pontes Reis, Eddy D Zandee van Rilland, Poonam Laxmappa Hosamani, Kevin R Keet, Minjoung Go, Evelyn Ling, David B Larson, Curtis Langlotz, Roxana Daneshjou, Jason Hom, Sanmi Koyejo, Emily Alsentzer, Akshay S Chaudhari.

*In submission*, 2025. [arXiv PDF](#).

## TECHNICAL SKILLS

---

**Languages:** Python, Java, MATLAB, C/C++, C#

**Libraries:** JAX, PyTorch, NumPy/SciPy, Pandas, Matplotlib, Keras, Ax, OpenAI, OpenCV

**Developer Tools:** Git, VS Code, Visual Studio

**Other:** machine learning, LLM, representation learning, causal inference, time-series analysis, L<sup>A</sup>T<sub>E</sub>X, Linux

## REFEREES

---

**Dr. Sanmi Koyejo**, Assistant Professor of CS  
*Stanford University*  
(217) 300-4907, [sanmi@cs.stanford.edu](mailto:sanmi@cs.stanford.edu)

**Dr. Eric Eaton**, Research Associate Professor of CS  
*University of Pennsylvania*  
(215) 746-1734, [eeaton@seas.upenn.edu](mailto:eeaton@seas.upenn.edu)

**Dr. Rahul Mangharam**, Associate Professor of EECS  
*University of Pennsylvania*  
(215) 573-3636, [rahulm@seas.upenn.edu](mailto:rahulm@seas.upenn.edu)

**Dr. David Meaney**, Senior Associate Dean  
*University of Pennsylvania*  
(215) 573-3155, [dmeaney@seas.upenn.edu](mailto:dmeaney@seas.upenn.edu)

## OTHER INTERESTS

---

Hair stylist/colorist for the bold, unadventurous, and everyone in between. Obstacle course race athlete, dancer when nobody is watching, and towel animal extraordinaire. Baroque Ensemble harpsichordist and cellist. Piano soloist at Carnegie Hall who knows all the lyrics to *The Phantom of the Opera* by heart.